

Optimization and Machine Learning

Chih-Jen Lin
Department of Computer Science
National Taiwan University



Talk at Summer School on Optimization, Big Data and
Applications, July, 2017

Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Talk slides are available at http://www.csie.ntu.edu.tw/~cjlin/talks/italy_optml.pdf




Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions




What is Machine Learning?

- Extract knowledge from data and make predictions
- Representative tasks: classification, clustering, and others



Classification



Clustering

- We will focus on **classification**



Data Classification

- Given training data in different classes (labels known)
Predict test data (labels unknown)
- Classic example
 - Find a patient's blood pressure, weight, etc.
 - After several years, know if he/she recovers
 - Build a machine learning model
 - New patient: find blood pressure, weight, etc
 - Prediction
- Two main stages: training and testing



Why Is Optimization Used?

- Usually the goal of classification is to **minimize** the test error
- Therefore, many classification methods solve optimization problems



Optimization and Machine Learning

- Standard optimization packages may be directly applied to machine learning applications
- However, efficiency and scalability are issues
- Very often machine learning knowledge must be considered in designing suitable optimization methods



Optimization and Machine Learning (Cont'd)

- Sometimes optimization researchers fail to make real impact because of not knowing **the differences between the two areas**
- I like to talk about the connection between these two areas because I was trained as an optimization researcher but now work on machine learning



Optimization and Machine Learning (Cont'd)

- In this talk, I will discuss some lessons learned in developing two popular packages
 - LIBSVM: 2000–now
A library for support vector machines
 - LIBLINEAR: 2007–now
A library for large linear classification
- Let me shamelessly say a bit about how they have been widely used



Optimization and Machine Learning (Cont'd)

- LIBSVM is probably the most widely used SVM package. Its implementation paper has been cited more than **32,000** times (Google Scholar, 5/2017)
- LIBLINEAR is popularly used in Internet companies for large-scale linear classification



Minimizing Training Errors

- Basically a classification method starts with **minimizing the training errors**

$$\min_{\text{model}} \text{(training errors)}$$

- That is, all or most training data with labels should be correctly classified by our model
- A model can be a decision tree, a support vector machine, a neural network, or other types



Minimizing Training Errors (Cont'd)

- We consider the model to be a vector w
- That is, the decision function is

$$\text{sgn}(w^T x)$$


- For any data, x , the predicted label is

$$\begin{cases} 1 & \text{if } w^T x \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



Minimizing Training Errors (Cont'd)

- The two-dimensional situation



- This seems to be quite restricted, but practically x is in a much higher dimensional space



Minimizing Training Errors (Cont'd)

- To characterize the training error, we need a **loss function** $\xi(\mathbf{w}; \mathbf{x}, y)$ for each instance (\mathbf{x}, y)
- Ideally we should use 0–1 training loss:


$$\xi(\mathbf{w}; \mathbf{x}, y) = \begin{cases} 1 & \text{if } y\mathbf{w}^T \mathbf{x} < 0, \\ 0 & \text{otherwise} \end{cases}$$



Minimizing Training Errors (Cont'd)

- However, this function is **discontinuous**. The optimization problem becomes difficult

$$\xi(w; x, y)$$



- We need **continuous approximations**



Loss Functions

- Some commonly used ones:

$$\xi_{L1}(\mathbf{w}; \mathbf{x}, y) \equiv \max(0, 1 - y\mathbf{w}^T \mathbf{x}), \quad (1)$$


$$\xi_{L2}(\mathbf{w}; \mathbf{x}, y) \equiv \max(0, 1 - y\mathbf{w}^T \mathbf{x})^2, \quad (2)$$

$$\xi_{LR}(\mathbf{w}; \mathbf{x}, y) \equiv \log(1 + e^{-y\mathbf{w}^T \mathbf{x}}). \quad (3)$$

- SVM (Boser et al., 1992; Cortes and Vapnik, 1995): (1)-(2)
- Logistic regression (LR): (3)



Loss Functions (Cont'd)



Their performance is usually **similar**



Loss Functions (Cont'd)

These loss functions have **different differentiability**

ξ_{L1} : not differentiable

ξ_{L2} : differentiable but not twice differentiable

ξ_{LR} : twice differentiable

The same optimization method may **not** be applicable to all these losses



Loss Functions (Cont'd)

- However, minimizing training losses may not give a good model for future prediction
- Overfitting occurs







Overfitting

- See the illustration in the next slide
- For classification,
You can easily achieve 100% training accuracy
- This is useless
- When training a data set, we should
 - Avoid **underfitting**: small training error
 - Avoid **overfitting**: small testing error



● and ▲: training; ○ and △: testing



Regularization

- To minimize the training error we manipulate the w vector so that it fits the data
- To avoid overfitting we need a way to make w 's values **less extreme**.
- One idea is to make w values closer to zero
- We can add, for example,

$$\frac{\mathbf{w}^T \mathbf{w}}{2} \quad \text{or} \quad \|\mathbf{w}\|_1$$

to the function that is minimized



Regularized Linear Classification

- Training data $\{y_i, \mathbf{x}_i\}, \mathbf{x}_i \in R^n, i = 1, \dots, I, y_i = \pm 1$
- I : # of data, n : # of features

$$\min_{\mathbf{w}} f(\mathbf{w}), \quad f(\mathbf{w}) \equiv \frac{\mathbf{w}^T \mathbf{w}}{2} + C \sum_{i=1}^I \xi(\mathbf{w}; \mathbf{x}_i, y_i)$$

- $\mathbf{w}^T \mathbf{w}/2$: regularization term (we have no time to talk about L1 regularization here)
- $\xi(\mathbf{w}; \mathbf{x}, y)$: loss function: we hope $y\mathbf{w}^T \mathbf{x} > 0$
- C : regularization parameter



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Kernel Methods

- Kernel methods are a class of classification techniques where major operations are conducted by **kernel evaluations**
- A representative example is support vector machine (Boser et al., 1992; Cortes and Vapnik, 1995)



Support Vector Classification

- Training data $(\mathbf{x}_i, y_i), i = 1, \dots, l, \mathbf{x}_i \in R^n, y_i = \pm 1$
- Minimizing training losses with regularization

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^l \max(1 - y_i(\mathbf{w}^T \phi(\mathbf{x}_i) + b), 0)$$

- Note that here we add a bias term b so the decision function is

$$\text{sgn}(\mathbf{w}^T \phi(\mathbf{x}) + b)$$



Support Vector Classification (Cont'd)

- Then the hyperplane does not pass through 0
- If n (# of features) is small, b may be important. Otherwise, it's not
- There are also historical reasons
- In our discussion we sometimes include b but sometimes do not



Mapping Data to a Higher Dimensional Space

- To better separate the data, we map data to a higher dimensional space

$$\phi(\mathbf{x}) = [\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots]^T.$$

- For example,

$$\frac{\text{weight}}{\text{height}^2}$$

is a useful new feature to check if a person overweights or not



Difficulties After Mapping Data to a High-dimensional Space

- # variables in w = dimensions of $\phi(x)$
- Infinite variables if $\phi(x)$ is infinite dimensional
- Cannot do an infinite-dimensional inner product for predicting a test instance

$$\text{sgn}(w^T \phi(x) + b)$$

- Use kernel trick to go back to a finite number of variables



Support Vector Classification (Cont'd)

- The **dual** problem (**finite** # variables)

$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \alpha^T Q \alpha - \mathbf{e}^T \alpha \\ \text{subject to} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, l \\ & \mathbf{y}^T \alpha = 0, \end{aligned}$$

where $Q_{ij} = y_i y_j \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ and $\mathbf{e} = [1, \dots, 1]^T$

- At optimum

$$\mathbf{w} = \sum_{i=1}^l \alpha_i y_i \phi(\mathbf{x}_i)$$

- Kernel: $K(\mathbf{x}_i, \mathbf{x}_j) \equiv \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$; closed form

Example: Gaussian (RBF) kernel: $e^{-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2}$



Kernel Tricks

- It can be shown that at optimum, \mathbf{w} is a **linear combination of training data**

$$\mathbf{w} = \sum_{i=1}^I y_i \alpha_i \phi(\mathbf{x}_i)$$

Proofs not provided here.

- Special $\phi(\mathbf{x})$** such that the decision function becomes

$$\begin{aligned}\text{sgn}(\mathbf{w}^T \phi(\mathbf{x}) + b) &= \text{sgn} \left(\sum_{i=1}^I y_i \alpha_i \phi(\mathbf{x}_i)^T \phi(\mathbf{x}) + b \right) \\ &= \text{sgn} \left(\sum_{i=1}^I y_i \alpha_i K(\mathbf{x}_i, \mathbf{x}) + b \right)\end{aligned}$$



Kernel Tricks (Cont'd)

- $\phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ needs a **closed** form
- Example: $\mathbf{x}_i \in R^3, \phi(\mathbf{x}_i) \in R^{10}$

$$\phi(\mathbf{x}_i) = [1, \sqrt{2}(x_i)_1, \sqrt{2}(x_i)_2, \sqrt{2}(x_i)_3, (x_i)_1^2, (x_i)_2^2, (x_i)_3^2, \sqrt{2}(x_i)_1(x_i)_2, \sqrt{2}(x_i)_1(x_i)_3, \sqrt{2}(x_i)_2(x_i)_3]^T$$

Then $\phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^2$.

- Kernel: $K(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^T \phi(\mathbf{y})$; common kernels:

$e^{-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2}$, (Radial Basis Function)

$(\mathbf{x}_i^T \mathbf{x}_j / a + b)^d$ (Polynomial kernel)



$K(x, y)$ can be inner product in **infinite** dimensional space. Assume $x \in R^1$ and $\gamma > 0$.

$$\begin{aligned}
 e^{-\gamma \|x_i - x_j\|^2} &= e^{-\gamma(x_i - x_j)^2} = e^{-\gamma x_i^2 + 2\gamma x_i x_j - \gamma x_j^2} \\
 &= e^{-\gamma x_i^2 - \gamma x_j^2} \left(1 + \frac{2\gamma x_i x_j}{1!} + \frac{(2\gamma x_i x_j)^2}{2!} + \frac{(2\gamma x_i x_j)^3}{3!} + \dots\right) \\
 &= e^{-\gamma x_i^2 - \gamma x_j^2} \left(1 \cdot 1 + \sqrt{\frac{2\gamma}{1!}} x_i \cdot \sqrt{\frac{2\gamma}{1!}} x_j + \sqrt{\frac{(2\gamma)^2}{2!}} x_i^2 \cdot \sqrt{\frac{(2\gamma)^2}{2!}} x_j^2 \right. \\
 &\quad \left. + \sqrt{\frac{(2\gamma)^3}{3!}} x_i^3 \cdot \sqrt{\frac{(2\gamma)^3}{3!}} x_j^3 + \dots\right) = \phi(x_i)^T \phi(x_j),
 \end{aligned}$$

where

$$\phi(x) = e^{-\gamma x^2} \left[1, \sqrt{\frac{2\gamma}{1!}} x, \sqrt{\frac{(2\gamma)^2}{2!}} x^2, \sqrt{\frac{(2\gamma)^3}{3!}} x^3, \dots\right]^T.$$




- We don't discuss the primal-dual relationship here
- In this lecture we focus more on optimization algorithms rather than optimization theory



Support Vector Classification (Cont'd)

Only x_i of $\alpha_i > 0$ used \Rightarrow support vectors



Large Dense Quadratic Programming

$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \alpha^T Q \alpha - \mathbf{e}^T \alpha \\ \text{subject to} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, l \\ & \mathbf{y}^T \alpha = 0 \end{aligned}$$

- $Q_{ij} \neq 0$, Q : an l by l **fully dense** matrix
- 50,000 training points: 50,000 variables:
 $(50,000^2 \times 8/2)$ bytes = 10GB RAM to store Q



Large Dense Quadratic Programming (Cont'd)

- Tradition optimization methods **cannot** be directly applied here because **Q cannot even be stored**
- Currently, decomposition methods (a type of coordinate descent methods) are commonly used in practice



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Decomposition Methods

- Working on some variables each time (e.g., Osuna et al., 1997; Joachims, 1998; Platt, 1998)
- Similar to coordinate-wise minimization
- Working set B , $N = \{1, \dots, I\} \setminus B$ fixed
- Let the objective function be

$$f(\alpha) = \frac{1}{2} \alpha^T Q \alpha - e^T \alpha$$



Decomposition Methods (Cont'd)

- Sub-problem on the variable \mathbf{d}_B

$$\begin{aligned} \min_{\mathbf{d}_B} \quad & f([\alpha_B] + [\mathbf{d}_B]) \\ \text{subject to} \quad & -\alpha_i \leq d_i \leq C - \alpha_i, \forall i \in B \\ & d_i = 0, \forall i \notin B, \\ & \mathbf{y}_B^T \mathbf{d}_B = 0 \end{aligned}$$

- The objective function of the sub-problem

$$\begin{aligned} & f([\alpha_B] + [\mathbf{d}_B]) \\ &= \frac{1}{2} \mathbf{d}_B^T Q_{BB} \mathbf{d}_B + \nabla_B f(\alpha)^T \mathbf{d}_B + \text{constant.} \end{aligned}$$



Avoid Memory Problems

- Q_{BB} is a sub-matrix of Q

$$\begin{bmatrix} Q_{BB} & Q_{BN} \\ Q_{NB} & Q_{NN} \end{bmatrix}$$

- Note that

$$\nabla f(\alpha) = Q\alpha - e, \quad \nabla_B f(\alpha) = Q_{B,:}\alpha - e_B$$



Avoid Memory Problems (Cont'd)

- Only B columns of Q are needed
- In general $|B| \leq 10$ is used. We need $|B| \geq 2$ because of the linear constraint

$$\mathbf{y}_B^T \mathbf{d}_B = 0$$

- Calculated when used: trade time for space
- But is such an approach practical?



Algorithm of Decomposition Methods

While α is not optimal

- Select a working set B
- Solve the sub-problem of d_B
- $\alpha_B \leftarrow \alpha_B + d_B$

We will talk about **the selection of B** later



How Decomposition Methods Perform?

- Convergence not very fast. This is known because of using only **first-order** information
- But, **no need** to have very accurate α decision function:

$$\text{sgn}(\mathbf{w}^T \phi(\mathbf{x}) + b) = \text{sgn} \left(\sum'_{i=1} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b \right)$$

Prediction may **still be correct** with a rough α

- Further, in some situations,
 - # support vectors \ll # training points
 - Initial $\alpha^1 = 0$, some instances **never used**



How Decomposition Methods Perform? (Cont'd)

- An example of training 50,000 instances using the software LIBSVM

```
$svm-train -c 16 -g 4 -m 400 22features  
Total nSV = 3370
```

Time 79.524s

- This was done on a typical desktop
 - Calculating the whole Q takes more time
 - #SVs = 3,370 \ll 50,000
- A good case where some remain at zero all the time



How Decomposition Methods Perform? (Cont'd)

- Because many $\alpha_i = 0$ in the end, we can develop a **shrinking** techniques

Variables are removed during the optimization procedure. **Smaller** problems are solved
- For example, if $\alpha_i = 0$ in more than 100 iterations, probably we can remove it
- More advanced techniques are possible
- Of course in the end we must check all variables again for the stopping condition



Machine Learning Properties are Useful in Designing Optimization Algorithms

We have seen that special properties of SVM contribute to the viability of decomposition methods

- For machine learning applications, no need to accurately solve the optimization problem
- Because some optimal $\alpha_i = 0$, decomposition methods may not need to update all the variables
- Also, we can use shrinking techniques to reduce the problem size during decomposition methods



Issues Related to Optimization

- Working set selection
- Asymptotic convergence
- Finite termination & stopping conditions
- Convergence rate
- Numerical issues



Working Set Selection

- In general we don't choose a large $|B|$ because of the following reasons
 - ① The sub-problem becomes expensive: $O(|B|^3)$
 - ② # iterations may not be significantly decreased
- Currently a popular setting is to choose $|B| = 2$

$$B = \{i, j\}$$

- This is called SMO (Sequential Minimal Optimization)



Working Set Selection (Cont'd)

- One idea is to use **gradient information**
- If

$$\alpha_i < C, y_i = 1, \text{ and } -\nabla_i f(\boldsymbol{\alpha}) > 0,$$

then we can enlarge α_i ;

- Therefore, one possibility is

$$i \in \arg \max \{-y_t \nabla f(\boldsymbol{\alpha})_t \mid \alpha_t < C, y_t = 1 \text{ or } \alpha_t > 0, y_t = -1\}$$

$$j \in \arg \min \{-y_t \nabla f(\boldsymbol{\alpha})_t \mid \alpha_t < C, y_t = -1 \text{ or } \alpha_t > 0, y_t = 1\}$$

- They somewhat correspond to the **maximal violation** of the optimality condition



Working Set Selection (Cont'd)

- This setting was first used in Joachims (1998)
- Can we use a more sophisticated one?
- It's difficult because cost is a concern
- For example, if we check the second-order information (e.g., objective-value reduction) of all $\{i, j\}$ pairs,

$\{i, j\}$ pairs,

then

$O(I^2)$

cost is needed



Working Set Selection (Cont'd)

- Currently in LIBSVM, we use

$$i \in \arg \max \{ -y_t \nabla f(\alpha)_t \mid \alpha_t < C, y_t = 1 \text{ or } \alpha_t > 0, y_t = -1 \}$$

and select j by second-order information (Fan et al., 2005)

- The cost is still

$$O(l)$$



Complexity of Decomposition Methods

Let's describe the algorithm again

While α is not optimal

- Select a working set B
- Solve the sub-problem of d_B
- $\alpha_B \leftarrow \alpha_B + d_B$



Complexity of Decomposition Methods (Cont'd)

- To construct the sub-problem,

$$\nabla_B f(\alpha) = Q_{B,:}\alpha - \mathbf{e}_B$$

needs

$$O(|B| \times I \times n)$$

if each kernel evaluation costs $O(n)$

- But for the working set selection, we need the **whole** $\nabla f(\alpha)$
- The cost can be as large as

$$O(I \times I \times n)$$



Complexity of Decomposition Methods (Cont'd)

- Fortunately, we can use

$$\begin{aligned}\nabla f(\alpha + \begin{bmatrix} \mathbf{d}_B \\ \mathbf{0} \end{bmatrix}) &= Q\alpha - \mathbf{e} + Q \begin{bmatrix} \mathbf{d}_B \\ \mathbf{0} \end{bmatrix} \\ &= \nabla f(\alpha) + Q_{:,B} \mathbf{d}_B\end{aligned}$$

- Note that $Q_{:,B}$ is available earlier
- Therefore, the cost of decomposition methods for kernel SVM is

$$O(|B| \times l \times n) \times \# \text{ iterations}$$



Differences between Optimization and Machine Learning

- The two topics may have different focuses. We give the following example
- The decomposition method we just discussed converges more slowly when C is large
- Using $C = 1$ on a data set
 - # iterations: 508
- Using $C = 5,000$
 - # iterations: 35,241



- Optimization researchers may rush to solve difficult cases of large C

That's what I did before

- It turns out that **large C is less used than small C**
- Recall that SVM solves

$$\frac{1}{2} \mathbf{w}^T \mathbf{w} + C(\text{sum of training losses})$$

- A large C means to **overfit** training data
- This does not give good test accuracy



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Linear and Kernel Classification

- We have
 - Kernel \Rightarrow map data to a higher space
 - Linear \Rightarrow use the original data; $\phi(\mathbf{x}) = \mathbf{x}$
- Intuitively, kernel should give superior accuracy than linear
- There are even theoretical proofs. Roughly speaking, from the Taylor expansion of the Gaussian (RBF) kernel

$$e^{-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2}$$

linear SVM is a **special case** of RBF-kernel SVM
(Keerthi and Lin, 2003)



Linear and Kernel Classification

- Optimization people may think there is no need to specially consider linear SVM
- That is, **same** optimization algorithms are enough for both linear and kernel cases
- However, this is wrong if we consider their practical use



Linear and Kernel Classification (Cont'd)

Classification methods such as SVM and logistic regression can be used in **two ways**

- Kernel methods: data mapped to a higher dimensional space

$$\mathbf{x} \Rightarrow \phi(\mathbf{x})$$

$\phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ easily calculated; **little control** on $\phi(\cdot)$

- Feature engineering** + linear classification :
We have \mathbf{x} without mapping. Alternatively, we can say that $\phi(\mathbf{x})$ is our \mathbf{x} ; **full control** on \mathbf{x} or $\phi(\mathbf{x})$



Linear and Kernel Classification (Cont'd)

- For some problems, **accuracy** by linear is as good as nonlinear
But **training and testing are much faster**
- This particularly happens for document classification
Number of features (bag-of-words model) very large
Data very **sparse** (i.e., few non-zeros)



Comparison Between Linear and Kernel (Training Time & Test Accuracy)

| Data set | #data | #features | Linear | | RBF Kernel | |
|--------------------|---------|-----------|--------|-------|------------|-------|
| | | | Time | Acc. | Time | Acc. |
| MNIST38 | 11,982 | 752 | 0.1 | 96.82 | 38.1 | 99.70 |
| ijcnn1 | 49,990 | 22 | 1.6 | 91.81 | 26.8 | 98.69 |
| covtype_multiclass | 464,810 | 54 | 1.4 | 76.37 | 46,695.8 | 96.11 |
| news20 | 15,997 | 1,355,191 | 1.1 | 96.95 | 383.2 | 96.90 |
| real-sim | 57,848 | 20,958 | 0.3 | 97.44 | 938.3 | 97.82 |
| yahoo-japan | 140,963 | 832,026 | 3.1 | 92.63 | 20,955.2 | 93.31 |
| webspam | 280,000 | 254 | 25.7 | 93.35 | 15,681.8 | 99.26 |



Comparison Between Linear and Kernel (Training Time & Test Accuracy)

| Data set | #data | #features | Linear | | RBF Kernel | |
|--------------------|---------|-----------|--------|-------|------------|-------|
| | | | Time | Acc. | Time | Acc. |
| MNIST38 | 11,982 | 752 | 0.1 | 96.82 | 38.1 | 99.70 |
| ijcnn1 | 49,990 | 22 | 1.6 | 91.81 | 26.8 | 98.69 |
| covtype_multiclass | 464,810 | 54 | 1.4 | 76.37 | 46,695.8 | 96.11 |
| news20 | 15,997 | 1,355,191 | 1.1 | 96.95 | 383.2 | 96.90 |
| real-sim | 57,848 | 20,958 | 0.3 | 97.44 | 938.3 | 97.82 |
| yahoo-japan | 140,963 | 832,026 | 3.1 | 92.63 | 20,955.2 | 93.31 |
| webspam | 280,000 | 254 | 25.7 | 93.35 | 15,681.8 | 99.26 |



Comparison Between Linear and Kernel (Training Time & Test Accuracy)

| Data set | #data | #features | Linear | | RBF Kernel | |
|--------------------|---------|-----------|--------|-------|------------|-------|
| | | | Time | Acc. | Time | Acc. |
| MNIST38 | 11,982 | 752 | 0.1 | 96.82 | 38.1 | 99.70 |
| ijcnn1 | 49,990 | 22 | 1.6 | 91.81 | 26.8 | 98.69 |
| covtype_multiclass | 464,810 | 54 | 1.4 | 76.37 | 46,695.8 | 96.11 |
| news20 | 15,997 | 1,355,191 | 1.1 | 96.95 | 383.2 | 96.90 |
| real-sim | 57,848 | 20,958 | 0.3 | 97.44 | 938.3 | 97.82 |
| yahoo-japan | 140,963 | 832,026 | 3.1 | 92.63 | 20,955.2 | 93.31 |
| webspam | 280,000 | 254 | 25.7 | 93.35 | 15,681.8 | 99.26 |

Therefore, there is a need to develop optimization methods for large linear classification



Why Linear is Faster in Training and Prediction?

- The reason is that
 - for linear, x_i is available
 - but
 - for kernel, $\phi(x_i)$ is not
- To illustrate this point, let's modify **kernel** decomposition methods discussed earlier for **linear**



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - **Decomposition method**
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Decomposition Method for Linear Classification

- Recall for kernels, we solve the sub-problem:

$$\min_{\mathbf{d}_B} \quad \frac{1}{2} \begin{bmatrix} (\boldsymbol{\alpha}_B + \mathbf{d}_B)^T & (\boldsymbol{\alpha}_N)^T \end{bmatrix} \begin{bmatrix} Q_{BB} & Q_{BN} \\ Q_{NB} & Q_{NN} \end{bmatrix} \begin{bmatrix} \boldsymbol{\alpha}_B + \mathbf{d}_B \\ \boldsymbol{\alpha}_N \end{bmatrix}$$

$$- [\mathbf{e}_B^T \quad \mathbf{e}_N^T] \begin{bmatrix} \boldsymbol{\alpha}_B + \mathbf{d}_B \\ \boldsymbol{\alpha}_N \end{bmatrix}$$

subject to $0 \leq \alpha_t \leq C, t \in B$, $\mathbf{y}_B^T(\boldsymbol{\alpha}_B + \mathbf{d}_B) = -\mathbf{y}_N^T \boldsymbol{\alpha}_N$

- The objective function over \mathbf{d}_B is

$$\frac{1}{2} \mathbf{d}_B^T Q_{BB} \mathbf{d}_B + (-\mathbf{e}_B + Q_{B,:} \boldsymbol{\alpha})^T \mathbf{d}_B$$



Decomposition Method for Linear Classification (Cont'd)

- We need

$$Q_{B,:}\boldsymbol{\alpha} = [Q_{BB} \quad Q_{BN}] \boldsymbol{\alpha}$$

- The cost is

$$O(|B| \times l \times n)$$

because for $i \in B$,

$$Q_{i,:}\boldsymbol{\alpha} = \sum_{j=1}^l y_i y_j K(\mathbf{x}_i, \mathbf{x}_j) \alpha_j$$

- n : # features, l : # data



Decomposition Method for Linear Classification (Cont'd)

- In the linear case,

$$K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$$

$$\Rightarrow \sum_{j=1}^I y_i y_j K(\mathbf{x}_i, \mathbf{x}_j) \alpha_j = y_i \mathbf{x}_i^T \left(\sum_{j=1}^I y_j \mathbf{x}_j \alpha_j \right)$$

- Assume

$$\mathbf{u} \equiv \sum_{j=1}^I y_j \alpha_j \mathbf{x}_j \quad (4)$$

is available



Decomposition Method for Linear Classification (Cont'd)

- The cost is significantly reduced

$$O(|B| \times l \times n) \Rightarrow O(|B| \times n)$$

- The main difference is that in kernel

$$\sum_{j=1}^l y_j \alpha_j \phi(\mathbf{x}_j)$$

cannot be written down. But for linear we can!



Decomposition Method for Linear Classification (Cont'd)

- All we need is to maintain \mathbf{u}

$$\mathbf{u} = \sum_{j=1}^J y_j \alpha_j \mathbf{x}_j$$

Then the following update rule can be applied

$$\mathbf{u} \leftarrow \mathbf{u} + \sum_{i:i \in B} (d_i) y_i \mathbf{x}_i.$$

- The cost is also

$$O(|B| \times n)$$



Decomposition Method for Linear Classification (Cont'd)

- Note that eventually

$\mathbf{u} \rightarrow$ primal optimal \mathbf{w}



One-variable Procedure

- This is what people use now in practice
- Let's consider the optimization problem without the bias term

$$\min_{\mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^I \max(1 - y_i \mathbf{w}^T \mathbf{x}_i, 0)$$

- The dual problem now **doesn't have a linear constraint**

$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \boldsymbol{\alpha}^T Q \boldsymbol{\alpha} - \mathbf{e}^T \boldsymbol{\alpha} \\ \text{subject to} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, I, \end{aligned}$$



One-variable Procedure (Cont'd)

- Here

$$f(\alpha) \equiv \frac{1}{2} \alpha^T Q \alpha - \mathbf{e}^T \alpha$$

and

$$Q_{ij} = y_i y_j \mathbf{x}_i^T \mathbf{x}_j, \quad \mathbf{e} = [1, \dots, 1]^T$$

- Without the linear constraint

$$\mathbf{y}^T \alpha = 0$$

we can choose one variable at a time



One-variable Procedure (Cont'd)

- Given current α . Let the working set be $B = \{i\}$
- Let

$$\mathbf{e}_i = [0, \dots, 0, 1, 0, \dots, 0]^T$$

- The sub-problem is

$$\min_d f(\alpha + d\mathbf{e}_i) = \frac{1}{2} Q_{ii} d^2 + \nabla_i f(\alpha) d + \text{constant}$$

subject to $0 \leq \alpha_i + d \leq C$

- Without constraints

$$\text{optimal } d = -\frac{\nabla_i f(\alpha)}{Q_{ii}}$$



One-variable Procedure (Cont'd)

- Now $0 \leq \alpha_i + d \leq C$

$$\alpha_i \leftarrow \min \left(\max \left(\alpha_i - \frac{\nabla_i f(\alpha)}{Q_{ii}}, 0 \right), C \right)$$

- Note that

$$\begin{aligned}\nabla_i f(\alpha) &= (Q\alpha)_i - 1 = \sum_{j=1}^I Q_{ij} \alpha_j - 1 \\ &= \sum_{j=1}^I y_i y_j \mathbf{x}_i^T \mathbf{x}_j \alpha_j - 1\end{aligned}$$



One-variable Procedure (Cont'd)

- As before, define

$$\mathbf{u} \equiv \sum_{j=1}^I y_j \alpha_j \mathbf{x}_j,$$

- Easy gradient calculation: the cost is $O(n)$

$$\nabla_i f(\boldsymbol{\alpha}) = y_i \mathbf{u}^T \mathbf{x}_i - 1$$



Algorithm: Dual Coordinate Descent

- Given initial α and find

$$\mathbf{u} = \sum_i y_i \alpha_i \mathbf{x}_i.$$

- While α is not optimal (Outer iteration)

For $i = 1, \dots, l$ (Inner iteration)

(a) $\bar{\alpha}_i \leftarrow \alpha_i$

(b) $G = y_i \mathbf{u}^T \mathbf{x}_i - 1$

(c) $\alpha_i \leftarrow \min(\max(\alpha_i - G/Q_{ii}, 0), C)$

(d) If α_i needs to be changed

$$\mathbf{u} \leftarrow \mathbf{u} + (\alpha_i - \bar{\alpha}_i) y_i \mathbf{x}_i$$



One-variable Procedure (Cont'd)

- Maintaining \mathbf{u} also costs

$$O(n)$$

- Thus the total cost is

$$O(n) \times \# \text{ iterations}$$

- Recall that the cost for **kernel** is

$$O(l \times n) \times \# \text{ iterations}$$

if we don't have b and select $|B| = 1$

- We will explain some interesting differences between the two



Recap: Dual Coordinate Descent

- It's very simple: minimizing one variable at a time
- While α not optimal

For $i = 1, \dots, I$

$$\min_{\alpha_i} f(\dots, \alpha_i, \dots)$$

- A classic optimization technique
- Traced back to Hildreth (1957) if constraints are not considered
- So what's new?



Recap: Dual Coordinate Descent (Cont'd)

- Having

$$\mathbf{u} \equiv \sum_{j=1}^I y_j \alpha_j \mathbf{x}_j,$$

$$\nabla_i f(\boldsymbol{\alpha}) = y_i \mathbf{u}^T \mathbf{x}_i - 1$$

and

$$\bar{\alpha}_i : \text{old} ; \quad \alpha_i : \text{new}$$

$$\mathbf{u} \leftarrow \mathbf{u} + (\alpha_i - \bar{\alpha}_i) y_i \mathbf{x}_i.$$

is very essential

- This is another example where we take the problem structure into account



Recap: Dual Coordinate Descent (Cont'd)

- Such a setting was first developed at Hsieh et al. (2008)



Careful Implementation

Some techniques can improve the running speed

- Shrinking: remove α_i if it is likely to be bounded at the end
Easier to conduct shrinking than the kernel case
(details not shown)
- Cyclic order to update elements

$$\alpha_1 \rightarrow \alpha_2 \rightarrow \cdots \rightarrow \alpha_I$$

A random order gives **faster convergence**

$$\alpha_{\pi(1)} \rightarrow \alpha_{\pi(2)} \rightarrow \cdots \rightarrow \alpha_{\pi(I)}$$

We will go back to this issue later



Difference from the Kernel Case

- We have seen that coordinate-descent type of methods are used for both linear and kernel classifiers
- Recall the i -th element of gradient costs $O(n)$ by

$$\begin{aligned}\nabla_i f(\alpha) &= \sum_{j=1}^I y_i y_j \mathbf{x}_i^T \mathbf{x}_j \alpha_j - 1 = (y_i \mathbf{x}_i)^T \left(\sum_{j=1}^I y_j \mathbf{x}_j \alpha_j \right) - 1 \\ &= (y_i \mathbf{x}_i)^T \mathbf{u} - 1\end{aligned}$$

but we **cannot** do this for kernel because

$$K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

cannot be separated



Difference from the Kernel Case (Cont'd)

- If using kernel, the cost of calculating $\nabla_i f(\alpha)$ must be $O(ln)$
- However, if $O(ln)$ cost is spent, the whole $\nabla f(\alpha)$ can be maintained

$$\nabla f(\alpha) \leftarrow \nabla f(\alpha) + Q_{:,i} d_i$$

- In contrast, the linear setting of using \mathbf{u} knows only $\nabla_i f(\alpha)$ rather than the whole $\nabla f(\alpha)$
- This situation affects the working set selection



Difference from the Kernel Case (Cont'd)

- We have mentioned that for existing decomposition methods (or say coordinate descent methods) for kernel classifiers, $\nabla f(\alpha)$ information is used to select variables for update
- Therefore, we have the following situations for two working set selections
 - ① Greedy: Using $\nabla f(\alpha)$
 - ② Cyclic

| | Kernel | | Linear | |
|-----------------------------|----------|----------|----------|--------|
| | Greedy | Cyclic | Greedy | Cyclic |
| Update α_i | $O(1)$ | $O(\ln)$ | $O(1)$ | $O(n)$ |
| Maintain $\nabla f(\alpha)$ | $O(\ln)$ | NA | $O(\ln)$ | NA |



Difference from the Kernel Case (Cont'd)

- In optimization there are two types of coordinate descent methods
 - ① Gauss-Seidel: **sequential** selection of variables
 - ② Gauss-Southwell: **greedy** selection of variables
- To do greedy selection, usually the whole gradient must be available
- Existing coordinate descent methods for linear \Rightarrow related to Gauss-Seidel
Existing coordinate descent methods for kernel \Rightarrow related to Gauss-Southwell



Difference from the Kernel Case (Cont'd)

- In general **greedy** leads to fewer iterations than **cyclic**
- So is the cyclic setting for linear practically viable?



Working Set Selection

- Without gradient information, looks like we can only do a cyclic update

1, 2, 3, ...

- However, Hsieh et al. (2008, Section 3.1) showed that with **random permutation**, the convergence is much faster
- Let's try an example by using LIBLINEAR
- real-sim is a data set with

72,309 instances and 20,958 features



Working Set Selection (Cont'd)

- With random permutation

```
$ ./train ~/Desktop/real-sim  
optimization finished, #iter = 13  
Objective value = -4396.370629
```

Here an iteration means to go through all instances once (though shrinking is applied)

- Without random permutation (cyclic)

```
$ ./train ~/Desktop/real-sim  
optimization finished, #iter = 326  
Objective value = -4391.706239
```



Working Set Selection (Cont'd)

- Here L1 loss is used
- Both approaches are under the same stopping condition
- Let's see the algorithm with random permutation



Working Set Selection (Cont'd)

- While α is not optimal (Outer iteration)

Randomly permute $1, \dots, I$

For $i = 1, \dots, I$ (Inner iteration)

(a) $\bar{\alpha}_i \leftarrow \alpha_i$

(b) $G = y_i \mathbf{u}^T \mathbf{x}_i - 1$

(c) $\alpha_i \leftarrow \min(\max(\alpha_i - G/Q_{ii}, 0), C)$

(d) If α_i needs to be changed

$$\mathbf{u} \leftarrow \mathbf{u} + (\alpha_i - \bar{\alpha}_i) y_i \mathbf{x}_i$$

- Note that it's not useful if we only randomly permute data once



Working Set Selection (Cont'd)

- Randomly permute $1, \dots, I$

While α is not optimal (Outer iteration)

For $i = 1, \dots, I$ (Inner iteration)

(a) $\bar{\alpha}_i \leftarrow \alpha_i$

(b) $G = y_i \mathbf{u}^T \mathbf{x}_i - 1$

(c) $\alpha_i \leftarrow \min(\max(\alpha_i - G/Q_{ii}, 0), C)$

(d) If α_i needs to be changed

$$\mathbf{u} \leftarrow \mathbf{u} + (\alpha_i - \bar{\alpha}_i)y_i \mathbf{x}_i$$

- On the other hand, random CD works:



Working Set Selection (Cont'd)

- While α is not optimal (Outer iteration)

Randomly select i

(a) $\bar{\alpha}_i \leftarrow \alpha_i$

(b) $G = y_i \mathbf{u}^T \mathbf{x}_i - 1$

(c) $\alpha_i \leftarrow \min(\max(\alpha_i - G/Q_{ii}, 0), C)$

(d) If α_i needs to be changed

$$\mathbf{u} \leftarrow \mathbf{u} + (\alpha_i - \bar{\alpha}_i)y_i \mathbf{x}_i$$

- Turns out it is easier to analyze the complexity of the random CD; see Shalev-Shwartz and Zhang (2013) and many subsequent works



Working Set Selection (Cont'd)

- It's difficult to analyze the setting of using random permutations. Some recent attempts (for simplified situations) include Lee and Wright (2016)
- This is still an ongoing research issue
- Another line of research is to **cheaply find some important indices.**
- See Glasmachers and Dogan (2013) and other developments
- In other words, we think some distributions are better than uniform



Working Set Selection (Cont'd)

- However, it is difficult to beat the uniform setting if shrinking has been applied



Working Set Selection and Bias Term

- Recall if we use

$$\text{sgn}(\mathbf{w}^T \mathbf{x} + b)$$

as the decision function, then the dual problem has a linear constraint

$$\mathbf{y}^T \boldsymbol{\alpha} = 0$$

- Then at each iteration, **two indices**

$$\{i, j\}$$

must be selected



Working Set Selection and Bias Term (Cont'd)

- If we use a cyclic setting, this means

$$(1, 2), (1, 3), (1, 4), \dots, (2, 3), \dots$$

- The training may be terribly slow because for many pairs we cannot change α at all
- Therefore, an interesting situation is as follows

| Working set selection | Without bias | With bias | Used for |
|-----------------------------------|--------------|-----------|----------|
| cyclic | OK | not OK | linear |
| greedy (with $\nabla f(\alpha)$) | OK | OK | kernel |



Working Set Selection and Bias Term (Cont'd)

- Fortunately, linear classification is often used to handle large data with many sparse features
- In such a high dimensional space, bias term is often not needed
- Even if it is, there is a trick of adding the bias term to the objective function

$$\min_{\mathbf{w}, b} \frac{1}{2} \mathbf{w}^T \mathbf{w} + \frac{1}{2} b^2 + C \sum_{i=1}^l \max(1 - y_i (\begin{bmatrix} \mathbf{w}^T & b \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}), 0)$$



Working Set Selection and Bias Term (Cont'd)

- The dual no longer has the linear constraint

$$\mathbf{y}^T \boldsymbol{\alpha} = 0$$

- However, for some problems (e.g., one-class SVM), a linear constraint must be there (details not shown)
- Then the training by coordinate descent for the linear case can be an issue



Other Losses

- Our discussion so far is for ℓ_1 loss
- All results can be applied to other losses such as ℓ_2 loss, logistic loss, etc.



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - **Newton methods**
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



More Optimization Methods can be Applied for Linear

- Recall that

$$\mathbf{w} = \sum_{i=1}^I y_i \alpha_i \phi(\mathbf{x}_i)$$

- Kernel: can **only** solve an optimization problem of α because \mathbf{w} is too high dimensional
- Linear: can solve **either w or α**
- We will show an example to minimize over \mathbf{w}



Newton Method

- Let's minimize a **twice-differentiable** function

$$\min_{\mathbf{w}} f(\mathbf{w})$$

- For example, logistic regression has

$$\min_{\mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^I \log \left(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i} \right).$$

- Newton direction at iterate \mathbf{w}^k

$$\min_s \quad \nabla f(\mathbf{w}^k)^T s + \frac{1}{2} s^T \nabla^2 f(\mathbf{w}^k) s$$



Truncated Newton Method

- The above sub-problem is equivalent to solving Newton linear system

$$\nabla^2 f(\mathbf{w}^k) \mathbf{s} = -\nabla f(\mathbf{w}^k)$$

- Approximately solving the linear system \Rightarrow truncated Newton
- However, Hessian matrix $\nabla^2 f(\mathbf{w}^k)$ is too large to be stored

$\nabla^2 f(\mathbf{w}^k) : n \times n, \quad n : \text{number of features}$

- For document data, n can be millions or more



Using Properties of Data Classification

- But Hessian has a special form

$$\nabla^2 f(\mathbf{w}) = \mathcal{I} + C \mathbf{X}^T \mathbf{D} \mathbf{X},$$

- \mathbf{D} is diagonal. For logistic regression,

$$D_{ii} = \frac{e^{-y_i \mathbf{w}^T \mathbf{x}_i}}{1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}}$$

- \mathbf{X} : data, # instances \times # features

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_l^T \end{bmatrix} \in R^{l \times n}$$



Using Properties of Data Classification (Cont'd)

- Using Conjugate Gradient (CG) to solve the linear system.
- CG is an iterative procedure. Each CG step mainly needs one Hessian-vector product

$$\nabla^2 f(\mathbf{w}) \mathbf{d} = \mathbf{d} + C \cdot X^T(D(X\mathbf{d}))$$

- Therefore, we have a Hessian-free approach



Using Properties of Data Classification (Cont'd)

- Now the procedure has **two layers** of iterations
 - Outer: Newton iterations
 - Inner: CG iterations per Newton iteration
- Past machine learning works used Hessian-free approaches include, for example, (Keerthi and DeCoste, 2005; Lin et al., 2008)
- Second-order information used: **faster convergence** than first-order methods



Training L2-loss SVM

- The loss function is differentiable but **not twice differentiable**

$$\xi_{L2}(\mathbf{w}; \mathbf{x}, y) \equiv \max(0, 1 - y\mathbf{w}^T \mathbf{x})^2$$

- We can use generalized Hessian (Mangasarian, 2002); details not shown here



Preconditioning

- Each Hessian-vector product

$$\nabla^2 f(\mathbf{w}) \mathbf{d} = \mathbf{d} + C \cdot X^T(D(X\mathbf{d}))$$

costs

$$O(ln), \quad \text{where } X \in R^{l \times n}$$

- Each function/gradient evaluation also costs $O(ln)$;
details omitted
- Therefore, the cost of each Newton iteration is
roughly

$$O(ln) \times \# \text{ CG steps}$$



Preconditioning (Cont'd)

- Can we reduce the number of CG steps by preconditioning?
- This classic technique finds

$$EE^T \approx \nabla^2 f(\mathbf{w})$$

so the linear system

$$\nabla^2 f(\mathbf{w})\mathbf{s} = -\nabla f(\mathbf{w})$$

becomes

$$E^{-1}\nabla^2 f(\mathbf{w})E^{-T}\hat{\mathbf{s}} = -E^{-1}\nabla f(\mathbf{w})$$



Preconditioning (Cont'd)

- Note that

$$s = E^{-T} \hat{s}$$

- If

$$E^{-1} \nabla^2 f(\mathbf{w}^k) E^{-T} \approx \mathcal{I}$$

then # CG steps can be reduced

- For example, we use the **diagonal preconditioner** so

$$EE^T = \text{diag}(\nabla^2 f(\mathbf{w}))$$

and

$$E = \sqrt{\text{diag}(\nabla^2 f(\mathbf{w}))}$$



Preconditioning (Cont'd)

- Because

$$\nabla^2 f(\mathbf{w}) = \mathcal{I} + C\mathbf{X}^T\mathbf{D}\mathbf{X},$$

we have

$$\nabla_{j,j}^2 f(\mathbf{w}) = 1 + C \sum_{i=1}^I D_{ii} X_{ij}^2$$

- The cost of constructing the preconditioner is

$$O(I n)$$



Preconditioning (Cont'd)

- In each CG step, doing

$$(E^{-1} \text{ or } E^{-T}) \times \text{ a vector}$$

costs

$$O(n)$$

- Therefore, extra cost of diagonal preconditioning at each Newton iteration is

$$O(n) \times \# \text{ CG steps} + O(ln)$$

- This is acceptable in compared with what we already need

$$O(ln) \times \# \text{ CG steps}$$



Preconditioning (Cont'd)

- However, the issue is that the number of CG steps may not be decreased
- In the area of numerical analysis, preconditioning is **an art rather than a science**
- Further, because of using a Hessian-free approach, many existing preconditioners are not directly applicable
- This is still a research issue



Sub-sampled Newton Methods

- The optimization problem can be rewritten as

$$\min_{\mathbf{w}} \quad \frac{\mathbf{w}^T \mathbf{w}}{2Cl} + \frac{1}{l} \sum_{i=1}^l \max(1 - y_i \mathbf{w}^T \mathbf{x}_i, 0)$$

- The second term is indeed the **average loss**
- A **subset** of data should give **similar** gradient or Hessian!!
- This kind of settings has been explored in Byrd et al. (2011); Martens (2010)
- Clearly this is an example of taking **properties of machine learning problems**



Lesson Learned from Kernel to Linear

- We must know the practical machine-learning use in order to decide if new optimization algorithms are needed for certain problems



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions




Comparisons

L2-loss SVM is used


- DCDL2: Dual coordinate descent (Hsieh et al., 2008)
- DCDL2-S: DCDL2 with shrinking (Hsieh et al., 2008)
- PCD: Primal coordinate descent (Chang et al., 2008)
- TRON: Trust region Newton method (Lin et al., 2008)




Objective values (Time in Seconds)




news20



rcv1



yahoo-japan



yahoo-korea


Analysis

- Dual coordinate descents are very effective if # data and # features are both large
Useful for document classification
- Half million data in **a few seconds**
- However, it is **less effective** if
features small: should solve **primal**; or
large penalty parameter C ; problems are more ill-conditioned




An Example When # Features Small

- # instance: 32,561, # features: 123



Objective value



Accuracy

Careful Evaluation

- This is **very very important**
- Let me give you one real personal experience



Careful Evaluation (Cont'd)

- In Lin et al. (2008), to check if diagonal perconditioner works in Newton methods we give the following table of **total number of CG steps**

| Problem | CG | PCG |
|-------------|-----|-----|
| a9a | 567 | 263 |
| real-sim | 104 | 160 |
| news20 | 71 | 155 |
| citeseer | 113 | 115 |
| yahoo-japan | 278 | 326 |
| rcv1 | 225 | 280 |
| yahoo-korea | 779 | 736 |




Careful Evaluation (Cont'd)

- Clearly, PCG may not be better
- But this conclusion is misleading
- In this experiment,
a strict stopping condition is used



Careful Evaluation (Cont'd)

- Let's look at this figure for yahoo-japan



- You see four horizontal lines



Careful Evaluation (Cont'd)

- The second line corresponds to the default LIBLINEAR stopping condition
- The 4 lines are by using

$$10\epsilon, \quad \epsilon, \quad 0.1\epsilon, \quad 0.01\epsilon,$$

where ϵ is the default LIBLINEAR stopping tolerance

- Things below/above these 4 lines are not useful
- They are too loose or too strict for the machine learning task



Careful Evaluation (Cont'd)

- Another issue is which regularization parameter C to be used in presenting results?
- In the table, we use $C = 1$, the simplest choice
- In the figure, we use

$$C_{\text{Best}} = 0.5,$$

where C_{Best} is the value leading to the **best** validation accuracy.

- A reasonable setting is to show figures of

$$C = C_{\text{Best}} \times \{0.01, 0.1, 1, 10, 100\},$$



Careful Evaluation (Cont'd)

- The reason is that in practice we start from a small C to search for C_{Best} .
- Things are larger than $100 \times C_{\text{Best}}$ are not important
- Lesson: even in comparing optimization algorithms for machine learning, we need to take machine learning properties into account



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Multi-core Implementation

- Nowadays each computer has several cores
- They can do parallel computation
- However, most machine learning algorithms (including those we have discussed) do not take parallel computation into account



Multi-core Implementation (Cont'd)

- In fact, algorithms may need to be redesigned
- Recently we did two studies
 - ① Newton method for solving the primal problem
(Lee et al., 2015)
 - ② Coordinate descent method for solving the dual problem
(Chiang et al., 2016)
- We will discuss the Newton method in detail



Parallel Newton Implementation

- Recall the bottleneck is the Hessian-vector product

$$\nabla^2 f(\mathbf{w}) \mathbf{d} = \mathbf{d} + C \cdot X^T(D(X\mathbf{d}))$$

- This is conducted at each CG step



Matrix-vector Multiplications: More Than 90% of the Training Time

| Data set | #instances | #features | ratio |
|--------------------|------------|------------|--------|
| kddb | 19,264,097 | 29,890,095 | 82.11% |
| url_combined | 2,396,130 | 3,231,961 | 94.83% |
| webspam | 350,000 | 16,609,143 | 97.95% |
| rcv1_binary | 677,399 | 47,236 | 97.88% |
| covtype_binary | 581,012 | 54 | 89.20% |
| epsilon_normalized | 400,000 | 2,000 | 99.88% |
| rcv1_multiclass | 518,571 | 47,236 | 97.04% |
| covtype_multiclass | 581,012 | 54 | 89.06% |



Matrix-vector Multiplications: More Than 90% of the Training Time (Cont'd)

- This result is by Newton methods using one core
- We should parallelize matrix-vector multiplications
- For $\nabla^2 f(\mathbf{w})\mathbf{d}$ we must calculate

$$\mathbf{u} = X\mathbf{d} \tag{5}$$

$$\mathbf{u} \leftarrow D\mathbf{u} \tag{6}$$

$$\bar{\mathbf{u}} = X^T \mathbf{u}, \text{ where } \mathbf{u} = DX\mathbf{d} \tag{7}$$

- Because D is diagonal, (6) is easy



Matrix-vector Multiplications: More Than 90% of the Training Time (Cont'd)

- We will discuss the parallelization of (5) and (7)
- They are more complicated
- For our problems, X is large and sparse
- Efficient **parallel sparse** matrix-vector multiplications are still a research issue



Parallel $X\mathbf{d}$ Operation

- Assume that X is in a **row-oriented** sparse format

$$X = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_l^T \end{bmatrix} \text{ and } \mathbf{u} = X\mathbf{d} = \begin{bmatrix} \mathbf{x}_1^T \mathbf{d} \\ \vdots \\ \mathbf{x}_l^T \mathbf{d} \end{bmatrix}$$

- we have the following simple loop

1: **for** $i = 1, \dots, l$ **do**

2: $u_i = \mathbf{x}_i^T \mathbf{d}$

3: **end for**

- Because the l inner products are **independent**, we can easily parallelize the loop by, for example, **OpenMP**



Parallel $X^T \mathbf{u}$ Operation

- For the other matrix-vector multiplication

$$\bar{\mathbf{u}} = X^T \mathbf{u}, \text{ where } \mathbf{u} = D X \mathbf{d},$$

we have

$$\bar{\mathbf{u}} = u_1 \mathbf{x}_1 + \cdots + u_l \mathbf{x}_l.$$

- Because matrix X is **row-oriented**, accessing **columns** in X^T is much easier than rows
- We can use the following loop

1: **for** $i = 1, \dots, l$ **do**

2: $\bar{\mathbf{u}} \leftarrow \bar{\mathbf{u}} + u_i \mathbf{x}_i$

3: **end for**



Parallel $X^T \mathbf{u}$ Operation (Cont'd)

- There is no need to store a separate X^T
- However, it is possible that threads on $u_{i_1}x_{i_1}$ and $u_{i_2}x_{i_2}$ want to update the same component \bar{u}_s at the same time:

```
1: for  $i = 1, \dots, /$  do in parallel  
2:   for  $(x_i)_s \neq 0$  do  
3:      $\bar{u}_s \leftarrow \bar{u}_s + u_i(x_i)_s$   
4:   end for  
5: end for
```



Atomic Operations for Parallel $X^T \mathbf{u}$

- An atomic operation can avoid other threads to write \bar{u}_s at the same time.

```
1: for  $i = 1, \dots, l$  do in parallel  
2:   for  $(x_i)_s \neq 0$  do  
3:     atomic:  $\bar{u}_s \leftarrow \bar{u}_s + u_i(x_i)_s$   
4:   end for  
5: end for
```

- However, **waiting time** can be a serious problem



Reduce Operations for Parallel $X^T \mathbf{u}$

- Another method is using temporary dense arrays maintained by each thread, and summing up them in the end
- That is, store

$$\hat{\mathbf{u}}^p = \sum_i \{u_i \mathbf{x}_i \mid i \text{ run by thread } p\}$$


and then

$$\bar{\mathbf{u}} = \sum_p \hat{\mathbf{u}}^p$$




Atomic Operation: Almost No Speedup

- Reduce operations are **superior** to atomic operations



rcv1_binary



covtype_binary

- Subsequently we use the reduce operations




Existing Algorithms for Sparse Matrix-vector Product

- Instead of our direct implementation to parallelize loops, in the next slides we will consider two existing methods



Recursive Sparse Blocks (Martone, 2014)

- RSB (Recursive Sparse Blocks) is an effective format for fast parallel sparse matrix-vector multiplications
- It recursively partitions a matrix to be like the figure
- Locality of memory references improved, but the **construction** time is not negligible



Recursive Sparse Blocks (Cont'd)

- Parallel, efficient sparse matrix-vector operations
- Improve locality of memory references
- But the initial construction time is about 20 multiplications, which is not negligible in some cases
- We will show the result in the experiments



Intel MKL

- Intel Math Kernel Library (MKL) is a commercial library including optimized routines for linear algebra (Intel)
- It supports fast matrix-vector multiplications for different sparse formats.
- We consider the row-oriented format to store X .



Experiments

- Baseline: Single core version in LIBLINEAR 1.96. It **sequentially** run the following operations

$$\mathbf{u} = X\mathbf{d}$$


$$\mathbf{u} \leftarrow D\mathbf{u}$$

$$\bar{\mathbf{u}} = X^T \mathbf{u}, \text{ where } \mathbf{u} = DX\mathbf{d}$$


- OpenMP: Use OpenMP to parallelize loops
- MKL: Intel MKL version 11.2
- RSB: librsb version 1.2.0



Speedup of Xd : All are Excellent



More Difficult to Speed up $X^T \mathbf{u}$



Indeed it's not easy to have a multi-core implementation
that is

- ① simple, and
- ② reasonably efficient

Let me describe what we do in the end in multi-core
LIBLINEAR



Reducing Memory Access to Improve Speedup

- In computing

$$\mathbf{X}\mathbf{d} \text{ and } \mathbf{X}^T(\mathbf{D}\mathbf{X}\mathbf{d})$$

the data matrix is **accessed twice**

- We notice that these two operations can be **combined together**


$$\mathbf{X}^T\mathbf{D}\mathbf{X}\mathbf{d} = \sum_{i=1}^I \mathbf{x}_i D_{ii} \mathbf{x}_i^T \mathbf{d}$$

- We can parallelize one **single** loop by OpenMP




Reducing Memory Access to Improve Speedup (Cont'd)

- Better speedup as memory accesses are reduced



rcv1_binary



covtype_binary

- The number of operations is **the same**, but memory access dramatically affects the idle time of threads



Reducing Memory Access to Improve Speedup (Cont'd)

- Therefore, if we can efficiently do

$$\sum_{i=1}^I \mathbf{x}_i D_{ii} \mathbf{x}_i^T \mathbf{d} \quad (8)$$

then probably we don't need sophisticated sparse matrix-vector packages

- However, for a simple operation like (8) careful implementations are still needed



OpenMP Scheduling

- An OpenMP loop assigns tasks to different threads.
- The default `schedule(static)` splits indices to P blocks, where each contains I/P elements.
- However, as tasks may be **unbalanced**, we can have a dynamic scheduling – available threads are assigned to the next tasks.
- For example, `schedule(dynamic, 256)` implies that a thread works on 256 elements each time.
- Unfortunately, overheads occur for the dynamic task assignment.



OpenMP Scheduling (Cont'd)

- Deciding suitable scheduling is not trivial.
- Consider implementing $X^T \mathbf{u}$ as an example. This operation involves the following three loops.
 - Initializing $\hat{\mathbf{u}}^p = \mathbf{0}, \forall p = 1, \dots, P$.
 - Calculating $\hat{\mathbf{u}}^p, \forall p$ by

$$\hat{\mathbf{u}}^p = \sum \{ \mathbf{u}_i \mathbf{x}_i \mid i \text{ run by thread } p \}$$

- Calculating $\bar{\mathbf{u}} = \sum_{p=1}^P \hat{\mathbf{u}}^p$.



OpenMP Scheduling (Cont'd)


- Consider the second step

| | covtype_binary | rcv1_binary |
|------------------------|----------------|-------------|
| schedule(static) | 0.2879 | 2.9387 |
| schedule(dynamic) | 1.2611 | 2.6084 |
| schedule(dynamic, 256) | 0.2558 | 1.6505 |

- Clearly, a suitable scheduling is essential
- The other two steps are more balanced, so `schedule(static)` is used (details omitted)



Speedup of Total Training Time



Analysis of Experimental Results

- For RSB, the speedup for $X\mathbf{d}$ is excellent, but is poor for $X^T \mathbf{u}$ on some $n \ll l$ data (e.g. covtype)
Furthermore, **construction time** is expensive
- OpenMP is the **best** for almost all cases, mainly because of **combing** $X\mathbf{d}$ and $X^T \mathbf{u}$ together
- Therefore, with **appropriate** settings, **simple** implementations by OpenMP can achieve excellent speedup



Outline

- 1 Introduction: why optimization and machine learning are related?
- 2 Optimization methods for kernel support vector machines
 - Decomposition methods
- 3 Optimization methods for linear classification
 - Decomposition method
 - Newton methods
 - Experiments
- 4 Multi-core implementation
- 5 Discussion and conclusions



Conclusions

- Optimization has been very useful for machine learning
- We must incorporate machine learning knowledge in designing suitable optimization algorithms and software
- The interaction between optimization and machine learning is very interesting and exciting.



References I

- B. E. Boser, I. Guyon, and V. Vapnik. A training algorithm for optimal margin classifiers. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, pages 144–152. ACM Press, 1992.
- R. H. Byrd, G. M. Chin, W. Neveitt, and J. Nocedal. On the use of stochastic Hessian information in optimization methods for machine learning. *SIAM Journal on Optimization*, 21(3):977–995, 2011.
- K.-W. Chang, C.-J. Hsieh, and C.-J. Lin. Coordinate descent method for large-scale L2-loss linear SVM. *Journal of Machine Learning Research*, 9:1369–1398, 2008. URL <http://www.csie.ntu.edu.tw/~cjlin/papers/cd12.pdf>.
- W.-L. Chiang, M.-C. Lee, and C.-J. Lin. Parallel dual coordinate descent method for large-scale linear classification in multi-core environments. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, 2016. URL http://www.csie.ntu.edu.tw/~cjlin/papers/multicore_cddual.pdf.
- C. Cortes and V. Vapnik. Support-vector network. *Machine Learning*, 20:273–297, 1995.
- R.-E. Fan, P.-H. Chen, and C.-J. Lin. Working set selection using second order information for training SVM. *Journal of Machine Learning Research*, 6:1889–1918, 2005. URL <http://www.csie.ntu.edu.tw/~cjlin/papers/quadworkset.pdf>.



References II

- T. Glasmachers and U. Dogan. Accelerated coordinate descent with adaptive coordinate frequencies. In *Proceedings of the 5th Asian Conference on Machine Learning*, volume 29 of *Proceedings of Machine Learning Research*, pages 72–86, 2013.
- C. Hildreth. A quadratic programming procedure. *Naval Research Logistics Quarterly*, 4: 79–85, 1957.
- C.-J. Hsieh, K.-W. Chang, C.-J. Lin, S. S. Keerthi, and S. Sundararajan. A dual coordinate descent method for large-scale linear SVM. In *Proceedings of the Twenty Fifth International Conference on Machine Learning (ICML)*, 2008. URL <http://www.csie.ntu.edu.tw/~cjlin/papers/cddual.pdf>.
- Intel. *Intel Math Kernel Library Reference Manual*.
- T. Joachims. Making large-scale SVM learning practical. In B. Schölkopf, C. J. C. Burges, and A. J. Smola, editors, *Advances in Kernel Methods – Support Vector Learning*, pages 169–184, Cambridge, MA, 1998. MIT Press.
- S. S. Keerthi and D. DeCoste. A modified finite Newton method for fast solution of large scale linear SVMs. *Journal of Machine Learning Research*, 6:341–361, 2005.
- S. S. Keerthi and C.-J. Lin. Asymptotic behaviors of support vector machines with Gaussian kernel. *Neural Computation*, 15(7):1667–1689, 2003.
- C.-P. Lee and S. J. Wright. Random permutations fix a worst case for cyclic coordinate descent, 2016. arXiv preprint arXiv:1607.08320.



References III

- M.-C. Lee, W.-L. Chiang, and C.-J. Lin. Fast matrix-vector multiplications for large-scale logistic regression on shared-memory systems. In *Proceedings of the IEEE International Conference on Data Mining (ICDM)*, 2015. URL
http://www.csie.ntu.edu.tw/~cjlin/papers/multicore_liblinear_icdm.pdf.
- C.-J. Lin, R. C. Weng, and S. S. Keerthi. Trust region Newton method for large-scale logistic regression. *Journal of Machine Learning Research*, 9:627–650, 2008. URL
<http://www.csie.ntu.edu.tw/~cjlin/papers/logistic.pdf>.
- O. L. Mangasarian. A finite Newton method for classification. *Optimization Methods and Software*, 17(5):913–929, 2002.
- J. Martens. Deep learning via Hessian-free optimization. In *Proceedings of the 27th International Conference on Machine Learning (ICML)*, 2010.
- M. Martone. Efficient multithreaded untransposed, transposed or symmetric sparse matrix–vector multiplication with the recursive sparse blocks format. *Parallel Computing*, 40:251–270, 2014.
- E. Osuna, R. Freund, and F. Girosi. Training support vector machines: An application to face detection. In *Proceedings of IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 130–136, 1997.



References IV

- J. C. Platt. Fast training of support vector machines using sequential minimal optimization. In B. Schölkopf, C. J. C. Burges, and A. J. Smola, editors, *Advances in Kernel Methods - Support Vector Learning*, Cambridge, MA, 1998. MIT Press.
- S. Shalev-Shwartz and T. Zhang. Stochastic dual coordinate ascent methods for regularized loss minimization. *Journal of Machine Learning Research*, 14:567–599, 2013.

